

Class 6: R functions

Neha (PID: A17567541)

2024-01-25

R functions

Functions are how we get stuff done. We call functions to do everything useful in R.

One cool thing about R is that it makes writing your own functions comparatively easy. All functions in R have at least 3 things:

- A **name** (we get to pick this)
- **Input arguments** (the input to our function)
- **The body** (lines of code that do the work)

```
#funname<- function(input1, input2) {  
  #The body with R code  
#}
```

Let's write a first function to add two numbers:

```
x<-5  
y<-1  
x+y
```

```
[1] 6
```

```
addme<-function(x,y=100) {x+y}
```

```
addme(1,1)
```

```
[1] 2
```

```
addme(10)
```

```
[1] 110
```

Lab for Today

```
# Example input vectors to start with
student1 <- c(100, 100, 100, 100, 100, 100, 100, 90)
student2 <- c(100, NA, 90, 90, 90, 90, 97, 80)
student3 <- c(90, NA, NA, NA, NA, NA, NA, NA)
```

```
mean(student1)
```

```
[1] 98.75
```

```
mean(student2, na.rm = TRUE)
```

```
[1] 91
```

Is there a way to set the value of NA to 0?.

Come back to the NA problem. We want to drop the lowest score before getting `mean()`

How do I find the lowest (minimum) score?

```
min(student1)
```

```
[1] 90
```

I found the `which.min` function. Maybe this is more useful?

```
#Tells you which element has the mininum
which.min(student1)
```

```
[1] 8
```

It is the 8th element of the vector that has the lowest score. Can I remove this one?

```
#Find the lowest score and remove it
student1[-which.min(student1)]
```

```
[1] 100 100 100 100 100 100 100
```

```
mean(student1[-which.min(student1)])
```

```
[1] 100
```

[] is used for which values you are calling out of a vector

Use a common shortcut and use x as my input

```
x<-student1
mean(x[-which.min(x)])
```

```
[1] 100
```

We still have the problem of missing values. One idea is to replace NA values with 0

This does not work.

```
y<-c(1,2, NA, 4, 5)
#| y==NA
```

```
is.na(y)
```

```
[1] FALSE FALSE  TRUE FALSE FALSE
```

How can I remove the NA element from the vector. Use ! to flip logicals.

```
y[!is.na(y)]
```

```
[1] 1 2 4 5
```

```
y[is.na(y)]<-1000
```

```
x<-student3

x[is.na(x)]<-0
#Changes the NA values to 0
mean(x[-which.min(x)])
```

```
[1] 12.85714
```

```
#Gives the mean
```

Last step now that I have my working code snippet is to make my `grade()` function.

```
grade()-function(x,y=1) {x[is.na(x)]<-0 #Changes the NA values to 0 mean(x[-
which.min(x)]) #Gives the mean}
```

Q1. Write a function `grade()` to determine an overall grade from a vector of student homework assignment scores dropping the lowest single score. If a student misses a homework (i.e. has an NA value) this can be used as a score to be potentially dropped. Your final function should be adequately explained with code comments and be able to work on an example class gradebook such as this one in CSV format: “<https://tinyurl.com/gradeinput>” [3pts]

```
grade<-function(x) {
  x[is.na(x)]<-0
  mean(x[-which.min(x)])
}
grade(student1)
```

```
[1] 100
```

Now read the online gradebook

```
url<-"https://tinyurl.com/gradeinput"
gradebook<-read.csv(url, row.names=1)
#row.names function sets the names of the rows. In this case, row.names 1 will set it to c

head(gradebook)
```

	hw1	hw2	hw3	hw4	hw5
student-1	100	73	100	88	79
student-2	85	64	78	89	78

```
student-3 83 69 77 100 77
student-4 88 NA 73 100 76
student-5 88 100 75 86 79
student-6 89 78 100 89 77
```

```
results<-apply(gradebook,1,grade)
print(results)
```

```
student-1 student-2 student-3 student-4 student-5 student-6 student-7
91.75      82.50      84.25      84.25      88.25      89.00      94.00
student-8 student-9 student-10 student-11 student-12 student-13 student-14
93.75      87.75      79.00      86.00      91.75      92.25      87.75
student-15 student-16 student-17 student-18 student-19 student-20
78.75      89.50      88.00      94.50      82.75      82.75
```

Q2. Using your grade() function and the supplied gradebook, Who is the top scoring student overall in the gradebook? [3pts]

```
max(results)
```

```
[1] 94.5
```

```
which.max(results)
```

```
student-18
18
```

Q3. From your analysis of the gradebook, which homework was toughest on students (i.e. obtained the lowest scores overall)? [2pts]

```
grade1<-function(x) {
  x[is.na(x)]<-0
  mean(x)
}
#If we used the original grade function, it would subtract the lowest scores, which would
results2<-apply(gradebook,2,grade1)
results2
```

```
hw1 hw2 hw3 hw4 hw5
89.00 72.80 80.80 85.15 79.25
```

```
min(results2)
```

```
[1] 72.8
```

```
which.min(results2)
```

```
hw2
```

```
2
```

Or, we could add an additional argument to the `apply()` function

```
which.min (apply(gradebook, 2, mean, na.rm=T))
```

```
hw3
```

```
3
```

```
#na.rm removes the NA values
```

```
which.min (apply(gradebook, 2, sum, na.rm=T))
```

```
hw2
```

```
2
```

```
min(apply(gradebook, 2, sum, na.rm=T))
```

```
[1] 1456
```

```
(apply(gradebook, 2, sum, na.rm=T))
```

```
hw1 hw2 hw3 hw4 hw5
```

```
1780 1456 1616 1703 1585
```

Using the mean might give the wrong impression, because outliers may affect the data significantly. If we omit the NA values and compute the sum, we will get a better measure of how students did.

Q4. Optional Extension: From your analysis of the gradebook, which homework was most predictive of overall score (i.e. highest correlation with average grade score)? [1pt]

```
#Make all (or mask) NA to 0
mask<-gradebook
mask[is.na(mask)]<-0
mask
```

	hw1	hw2	hw3	hw4	hw5
student-1	100	73	100	88	79
student-2	85	64	78	89	78
student-3	83	69	77	100	77
student-4	88	0	73	100	76
student-5	88	100	75	86	79
student-6	89	78	100	89	77
student-7	89	100	74	87	100
student-8	89	100	76	86	100
student-9	86	100	77	88	77
student-10	89	72	79	0	76
student-11	82	66	78	84	100
student-12	100	70	75	92	100
student-13	89	100	76	100	80
student-14	85	100	77	89	76
student-15	85	65	76	89	0
student-16	92	100	74	89	77
student-17	88	63	100	86	78
student-18	91	0	100	87	100
student-19	91	68	75	86	79
student-20	91	68	76	88	76

We can use the `cor()` function for correlation analysis

```
cor(mask$hw2, results)
```

```
[1] 0.176778
```

```
cor(mask$hw3, results)
```

```
[1] 0.3042561
```

I need to use the `apply()` function to run this analysis over the whole course (i.e. masked gradebook)

```
apply(mask, 2, cor, results)
```

hw1	hw2	hw3	hw4	hw5
0.4250204	0.1767780	0.3042561	0.3810884	0.6325982